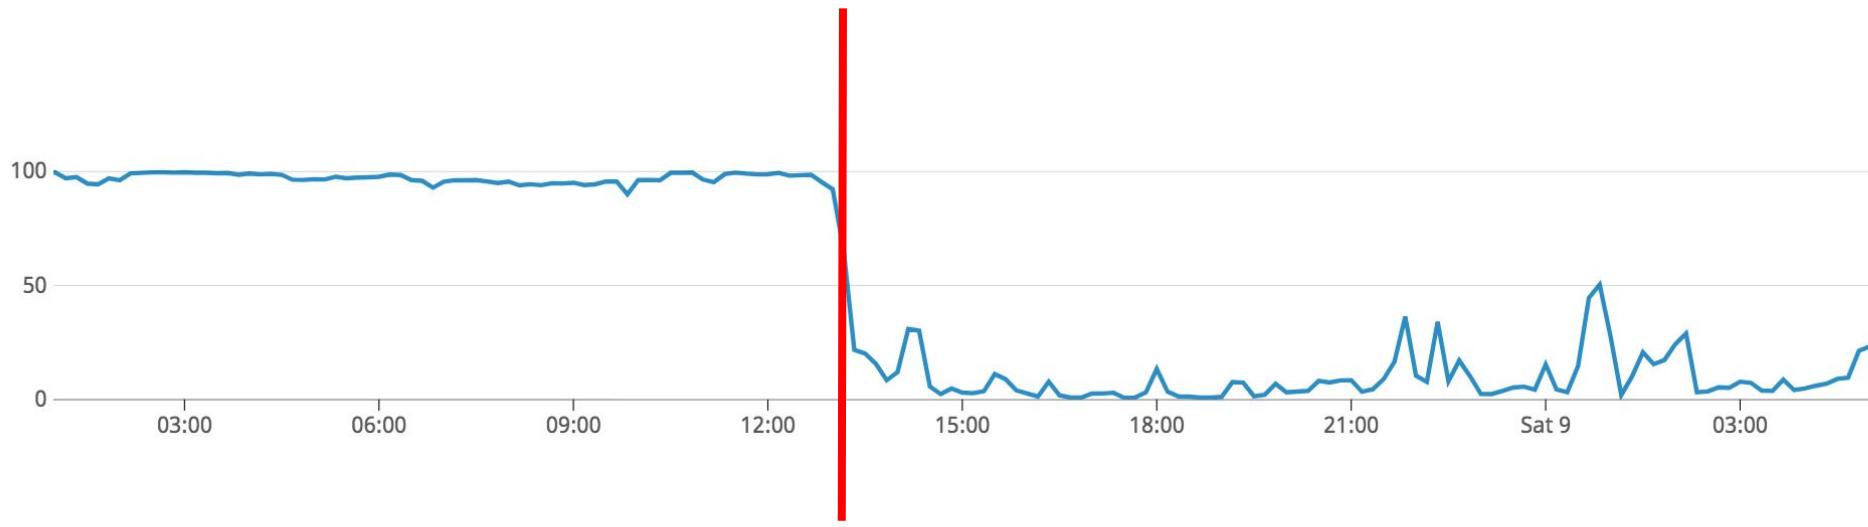




# Cache is King



@molly\_struve

KENNA  
Security



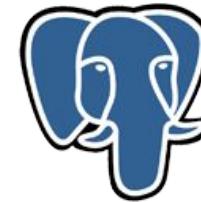
@molly\_struve





# Cache is King

# Site Reliability Engineer



PostgreSQL





## Adding Indexes

Using SELECT statements

Batch Processing

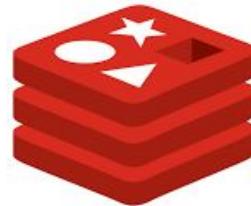


```
{  
  "statusCode": 200,  
  "took": "100ms"  
}
```

Elasticsearch::Tr

yTimeout 504

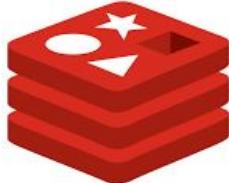
# Resque



redis



elastic

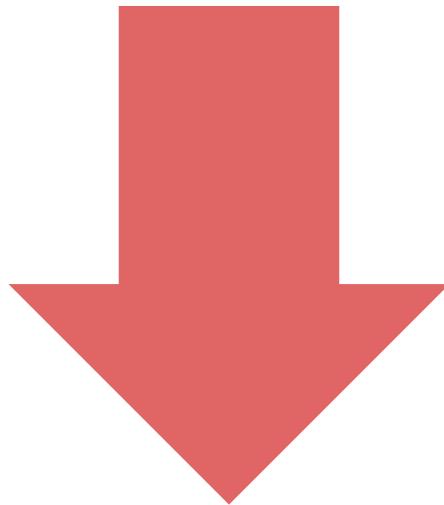


redis



# Demo Time!

# Quantity of Datastore Hits





# KENNA

## Security



HANES  
Brands Inc

@molly\_struve

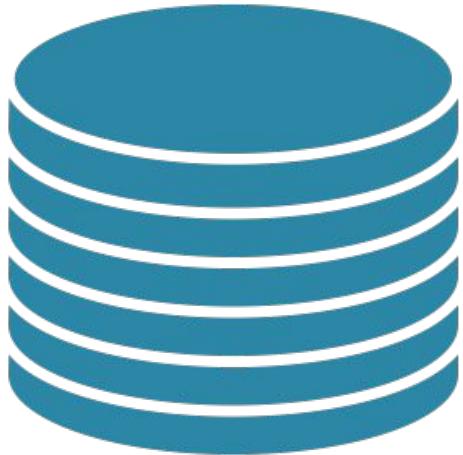
# The average company has...

60 thousand  
assets

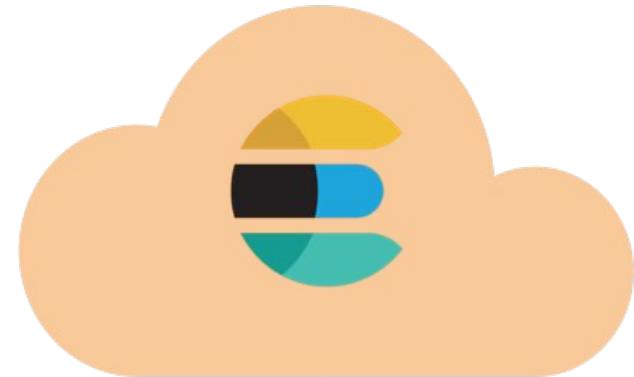
?

24 million  
vulnerabilities





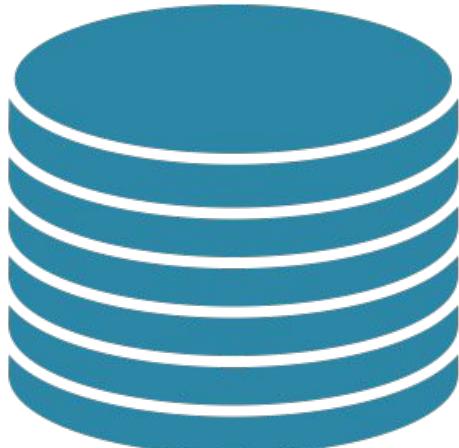
MySQL



Elasticsearch  
Cluster

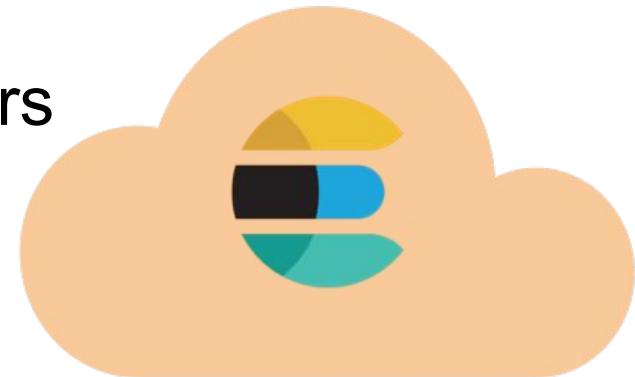
@molly\_struve

# Serialization



MySQL

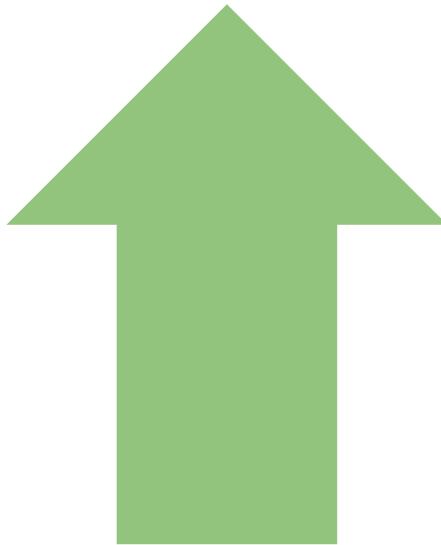
ActiveModelSerializers



Elasticsearch  
Cluster

@molly\_struve

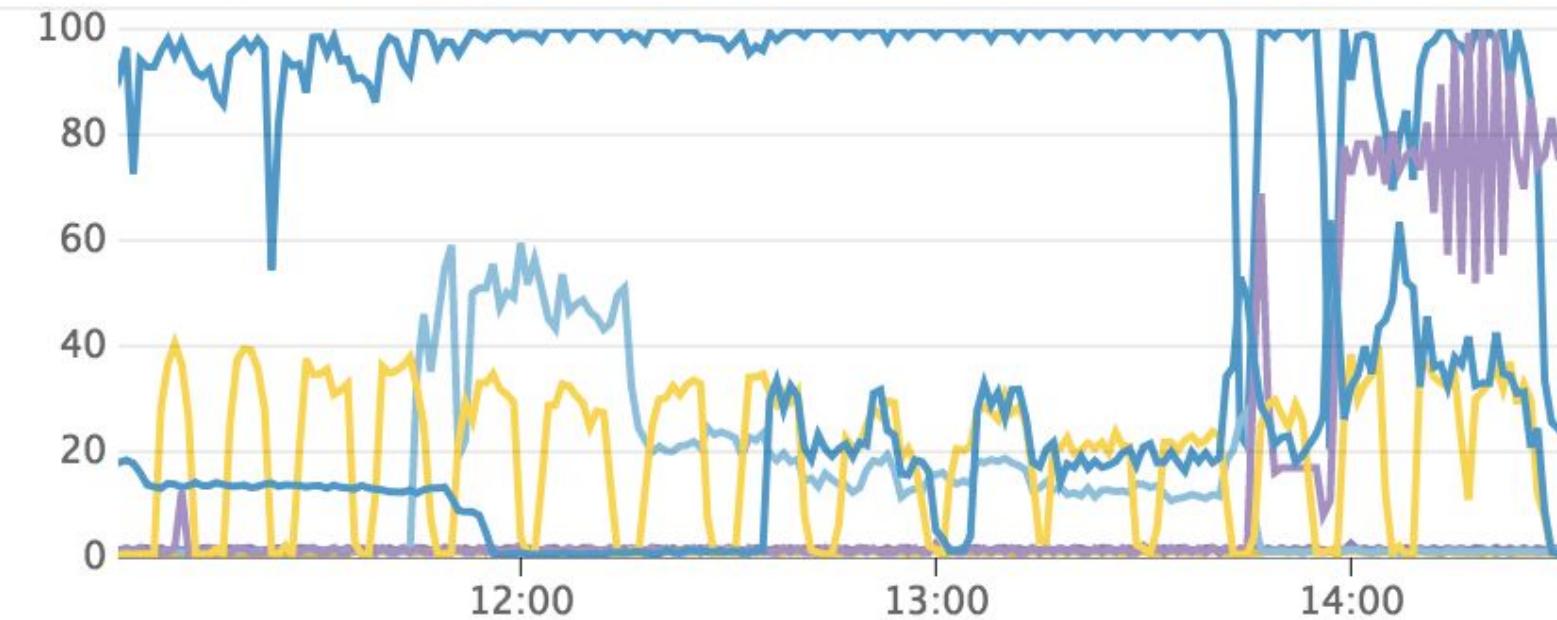
```
class Vulnerability < ActiveModel::Serializer
  attributes :id, :client_id, :created_at, :updated_at,
             :priority, :details, :notes, :asset_id,
             :solution_id, :owner_id, :ticket_id
end
```



**200 MILLION**

# 11 hours and counting...

RDS CPU utilization

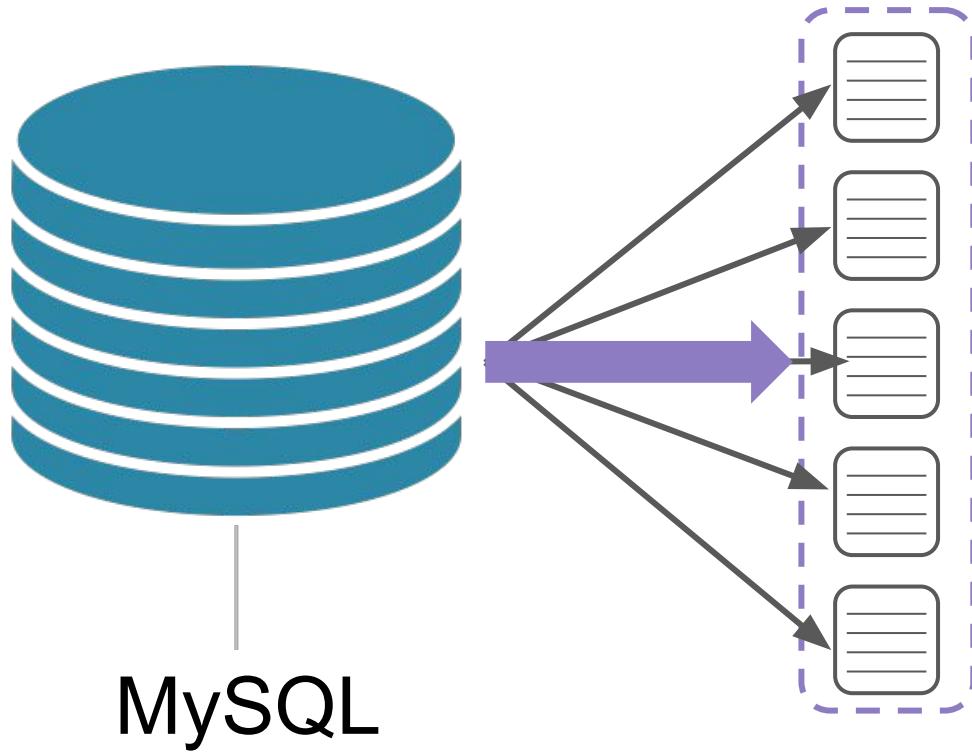


@molly\_struve





(2.2ms){0.9ms} (1.4ms)  
(2.9ms){6.0ms}  
(3.1ms) (2.1ms)  
(1.6ms){0.3ms}{5.2ms}{2.1ms}  
(0.9ms)  
(4.9ms){1.3ms}{1.3ms}{2.2ms}  
(5.2ms){3.0ms}{4.1ms}{6ms}



# Bulk Serialization

```
class BulkVulnerabilityCache
  attr_accessor :vulnerabilities, :client, :vulnerability_ids

  def initialize(vulns, client)
    self.vulnerabilities = vulns
    self.vulnerability_ids = vulns.map(&:id)
    self.client = client
  end

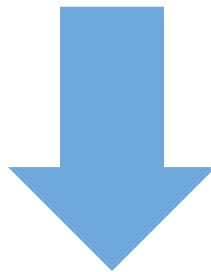
  # MySQL Lookups
end
```

```
module Serializers
  class Vulnerability
    attr_accessor :vulnerability, :cache

    def initialize(vuln, bulk_cache)
      self.cache = bulk_cache
      self.vulnerability = vuln
    end
  end
end
```

```
class Vulnerability
  has_many :custom_fields
end
```

CustomField.where(:vulnerability\_id => vuln.id)



cache.fetch('custom\_fields', vuln.id)

# The Result...

```
(pry)> vulns = Vulnerability.limit(300);
(pry)> Benchmark.realtime { vulns.each(&:serialize) }
=> 6.022452222998254

(pry)> Benchmark.realtime do
>   BulkVulnerability.new(vulns, [], client).serialize
> end
=> 0.7267019419959979
```



@molly\_struve

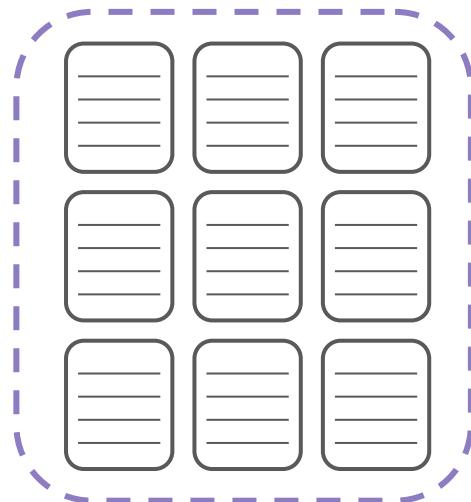
# Decrease in database hits

Individual Serialization: **2,100**

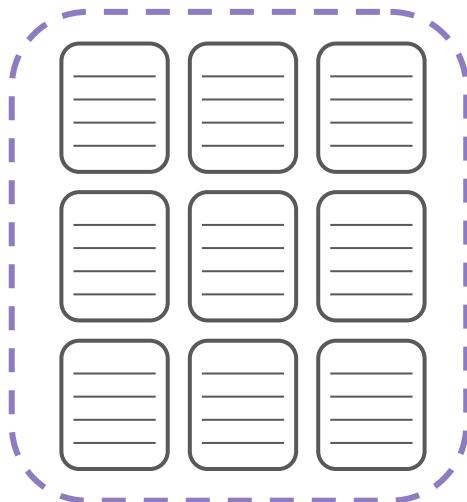
Bulk Serialization: **7**



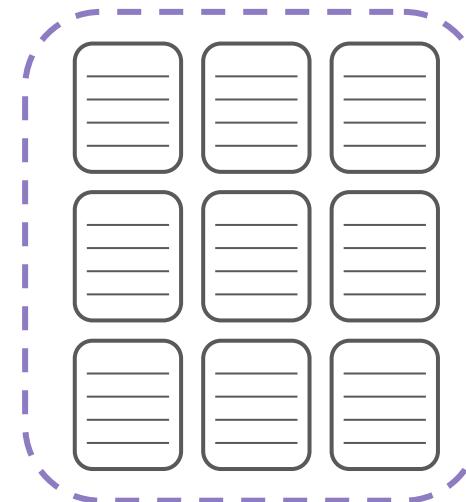
# Vulnerability Batches



**1k vulns**

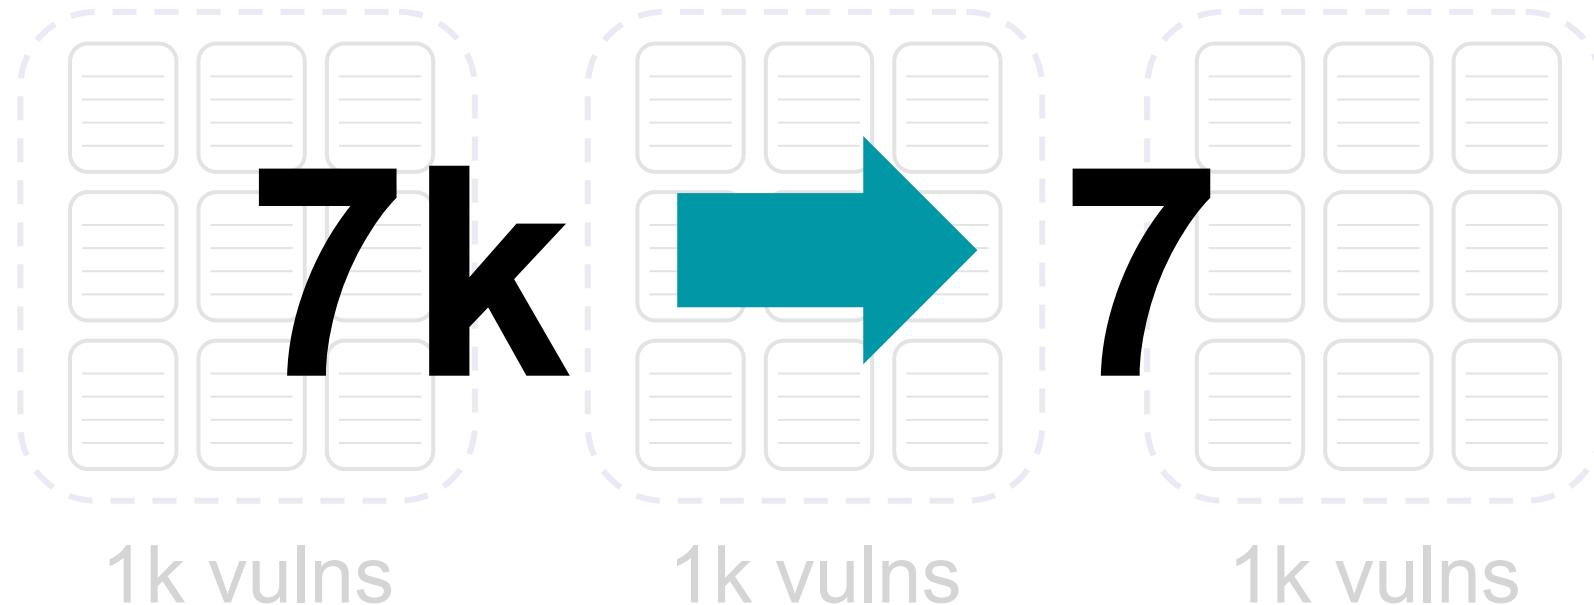


**1k vulns**

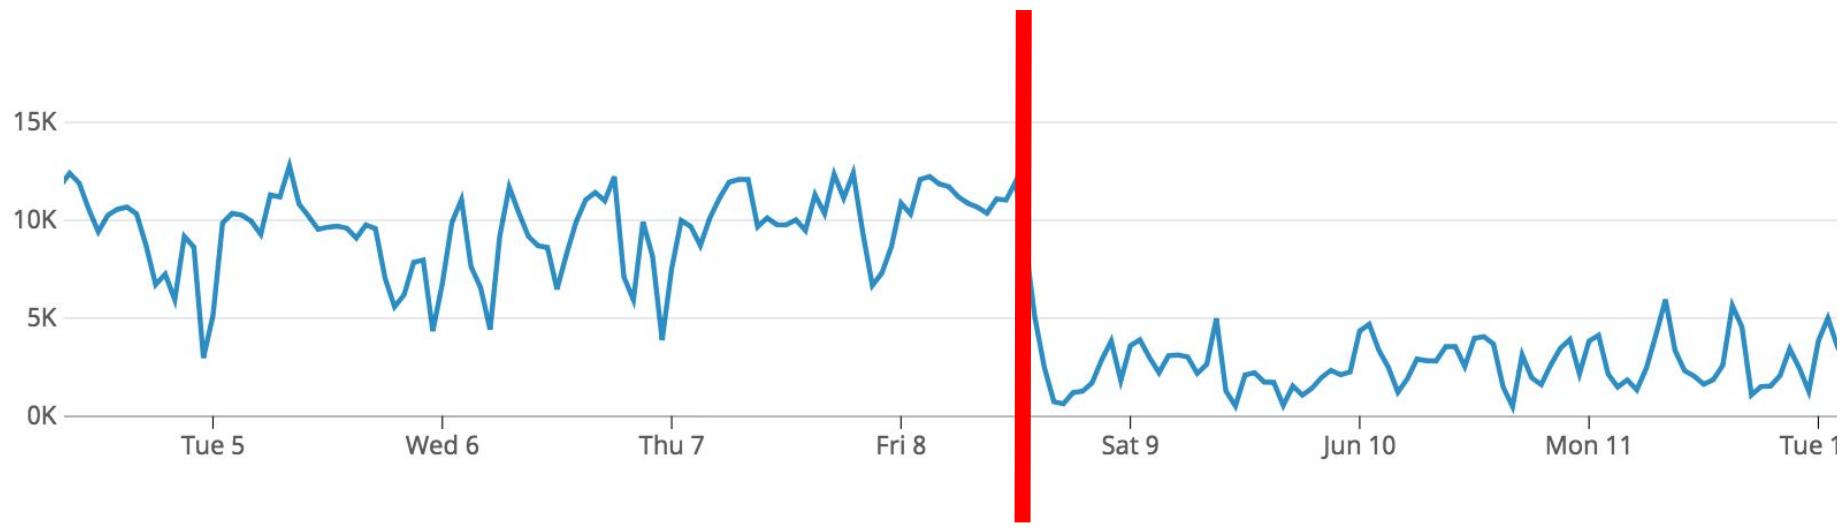


**1k vulns**

# Vulnerability Batches



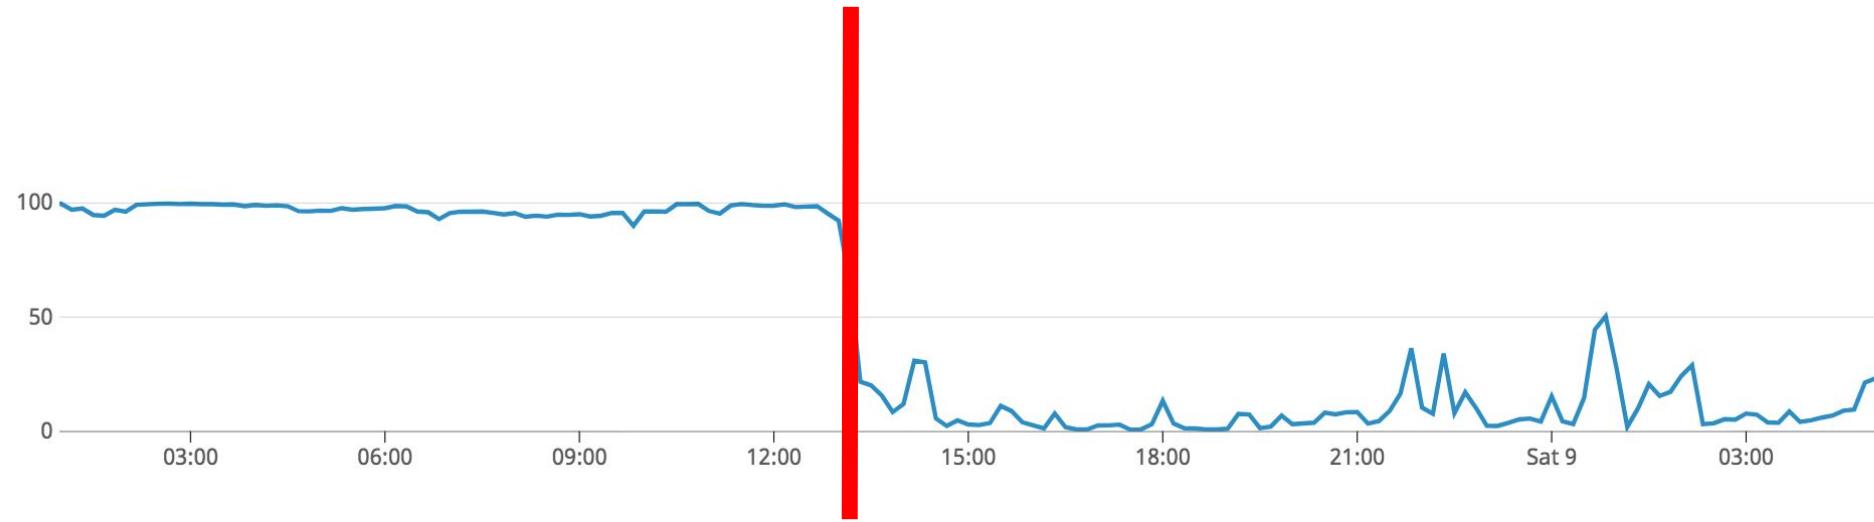
# MySQL Queries



Bulk Serialization  
Deployed

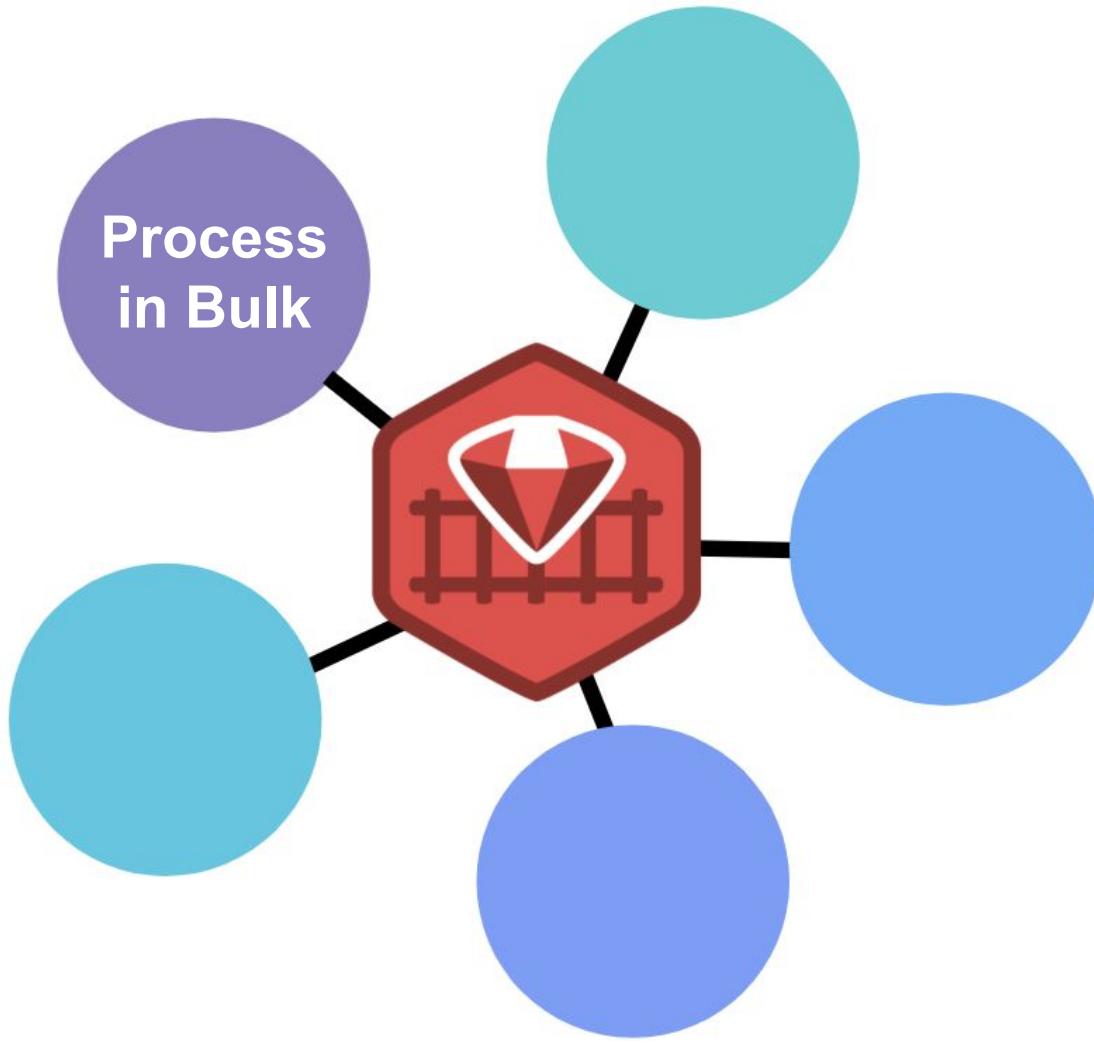
@molly\_struve

# RDS CPU Utilization



Bulk Serialization  
Deployed

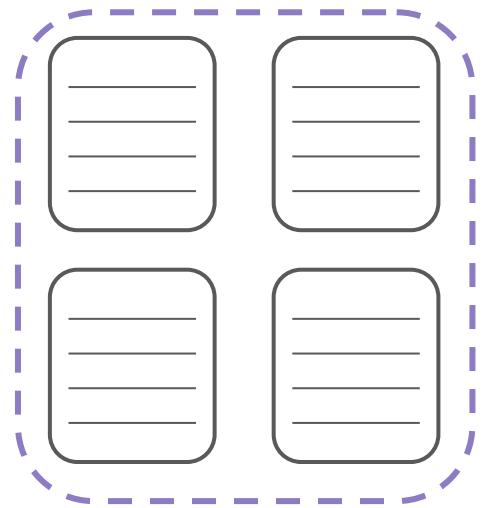
Process  
in Bulk







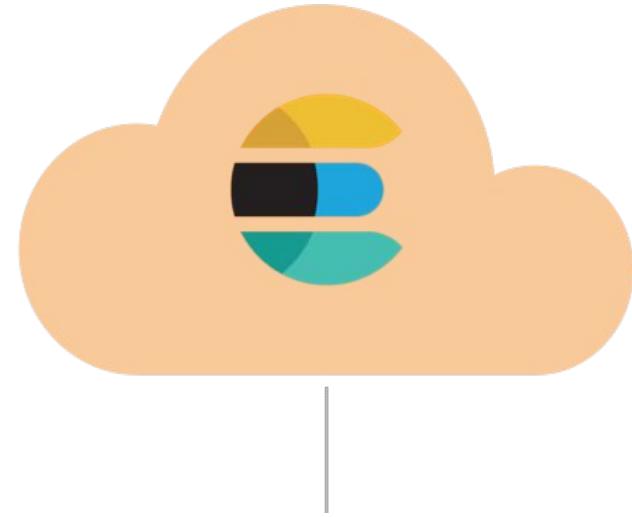
@molly\_struve



Vulnerabilities

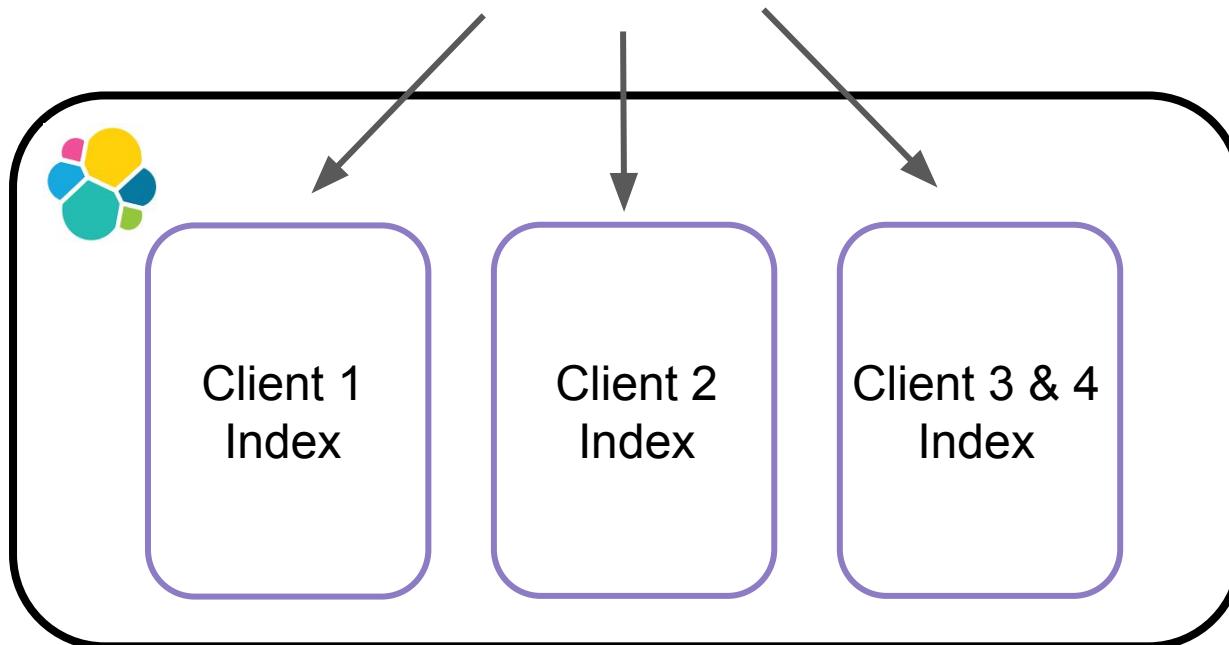


MySQL  
+  
Redis



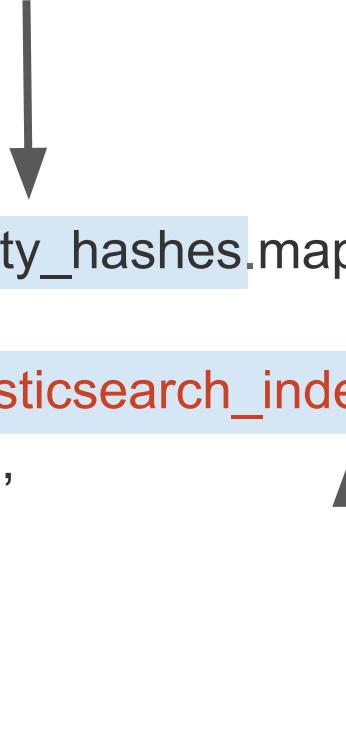
Elasticsearch  
Cluster

# Redis.get



```
indexing_hashes = vulnerability_hashes.map do |hash|
  {
    :_index => Redis.get("elasticsearch_index_#{hash[:client_id]}"),
    :_type => hash[:doc_type],
    :_id => hash[:id],
    :data => hash[:data]
  }
end
```

```
indexing_hashes = vulnerability_hashes.map do |hash|
  {
    :_index => Redis.get("elasticsearch_index_#{hash[:client_id]}"),
    :_type => hash[:doc_type],
    :_id => hash[:id],
    :data => hash[:data]
  }
end
```



# GET

```
(pry)> index_name = Redis.get("elasticsearch_index_#{client_id}")  
DEBUG -- : [Redis] command=GET args="elasticsearch_index_1234"  
DEBUG -- : [Redis] call_time=1.07 ms
```



# BulkIndexVulnFieldsWorker/perform



App performance

Map Beta

App server breakdown

10k ms

7.5k ms

5k ms

2500 ms

0

N/A  
APDEX

7.78 sec  
AVERAGE

Mar 20,  
3:25 PM

Mar 20,  
3:40 PM

Mar 20,  
3:55 PM

Mar 20,  
4:10 PM

Mar 20,  
4:25 PM

Breakdown table

Category

Segment

% Time

Avg calls  
(per txn)

Avg  
time (ms)

Database

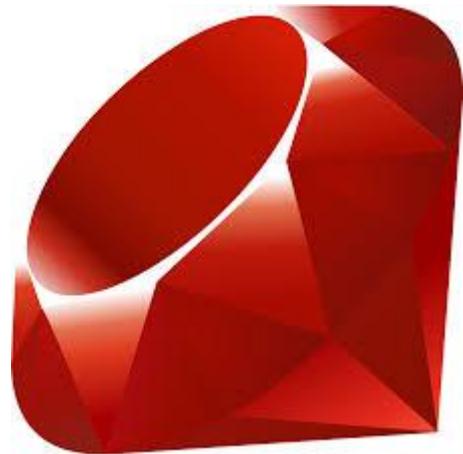
Redis get

64.9

4000

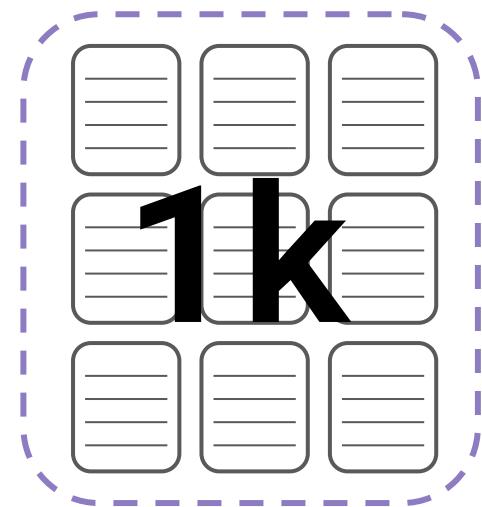
5,050

@molly\_struve



```
client_indexes = Hash.new do |h, client_id|
  h[client_id] = Redis.get("elasticsearch_index_#{client_id}")
end
```

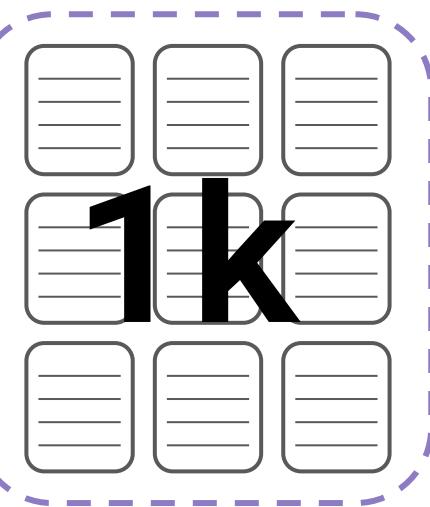
```
indexing_hashes = vuln_hashes.map do |hash|
{
  _index => client_indexes[hash[:client_id]],
  _type => hash[:doc_type],
  _id => hash[:id],
  :data => hash[:data]
}
end
```



Client 1

1

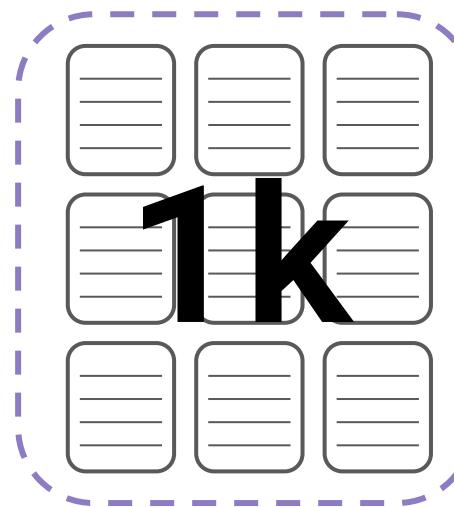
+



Client 2

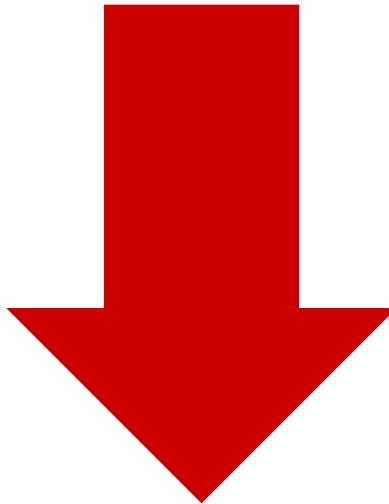
1

+

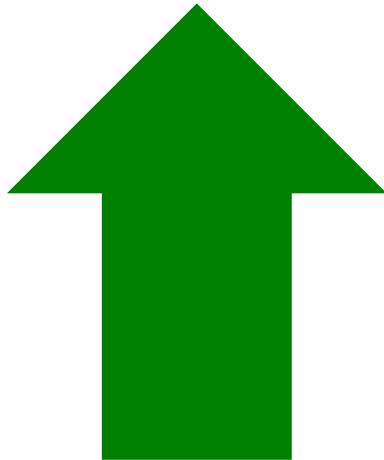


Client 3

1



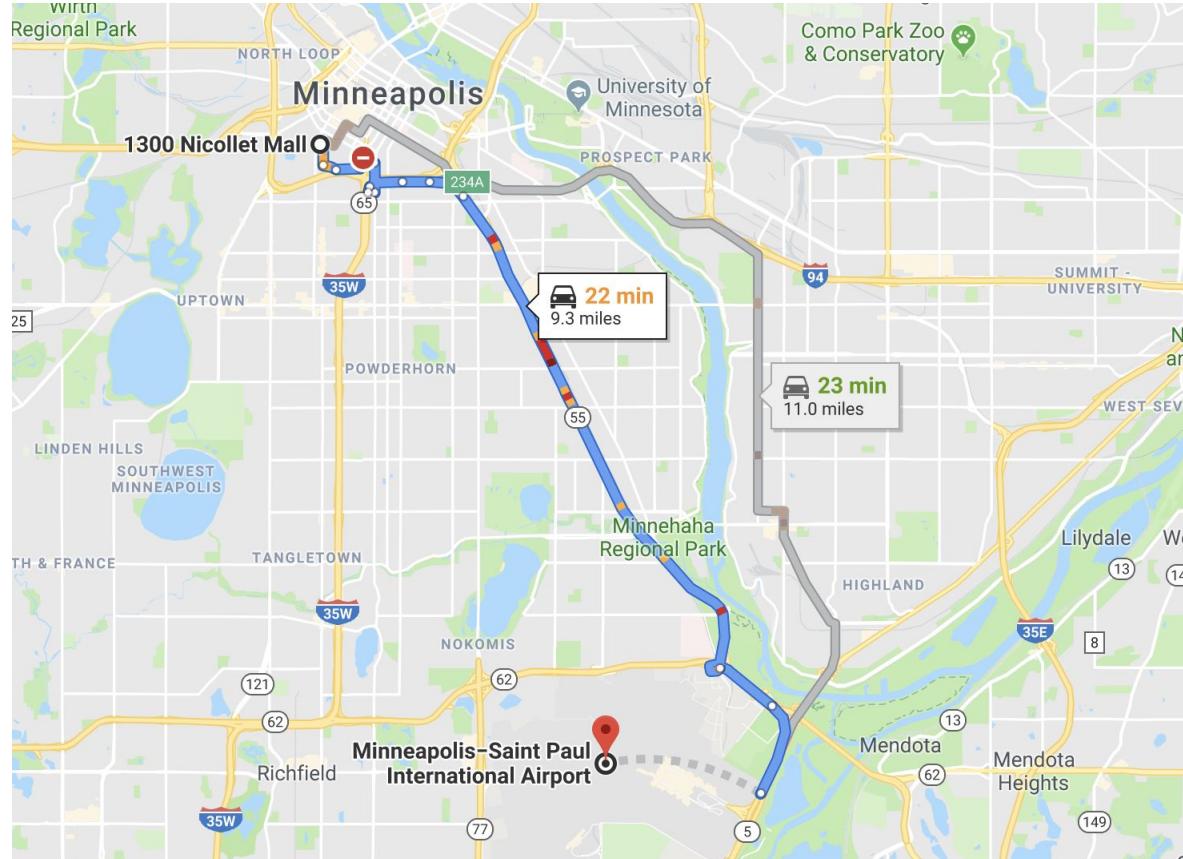
1 000x



**65% job  
speed up**

# Local Cache

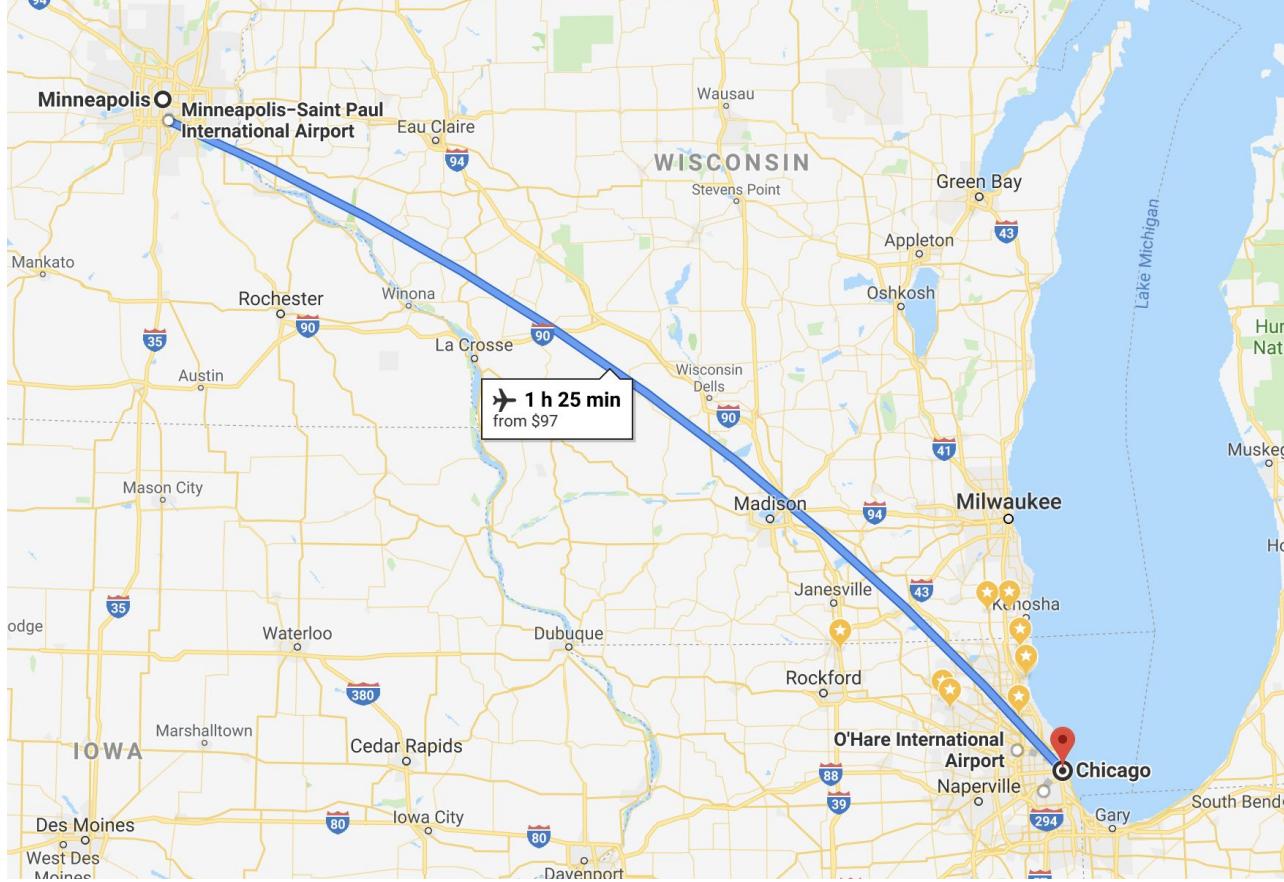
KENNA  
Security



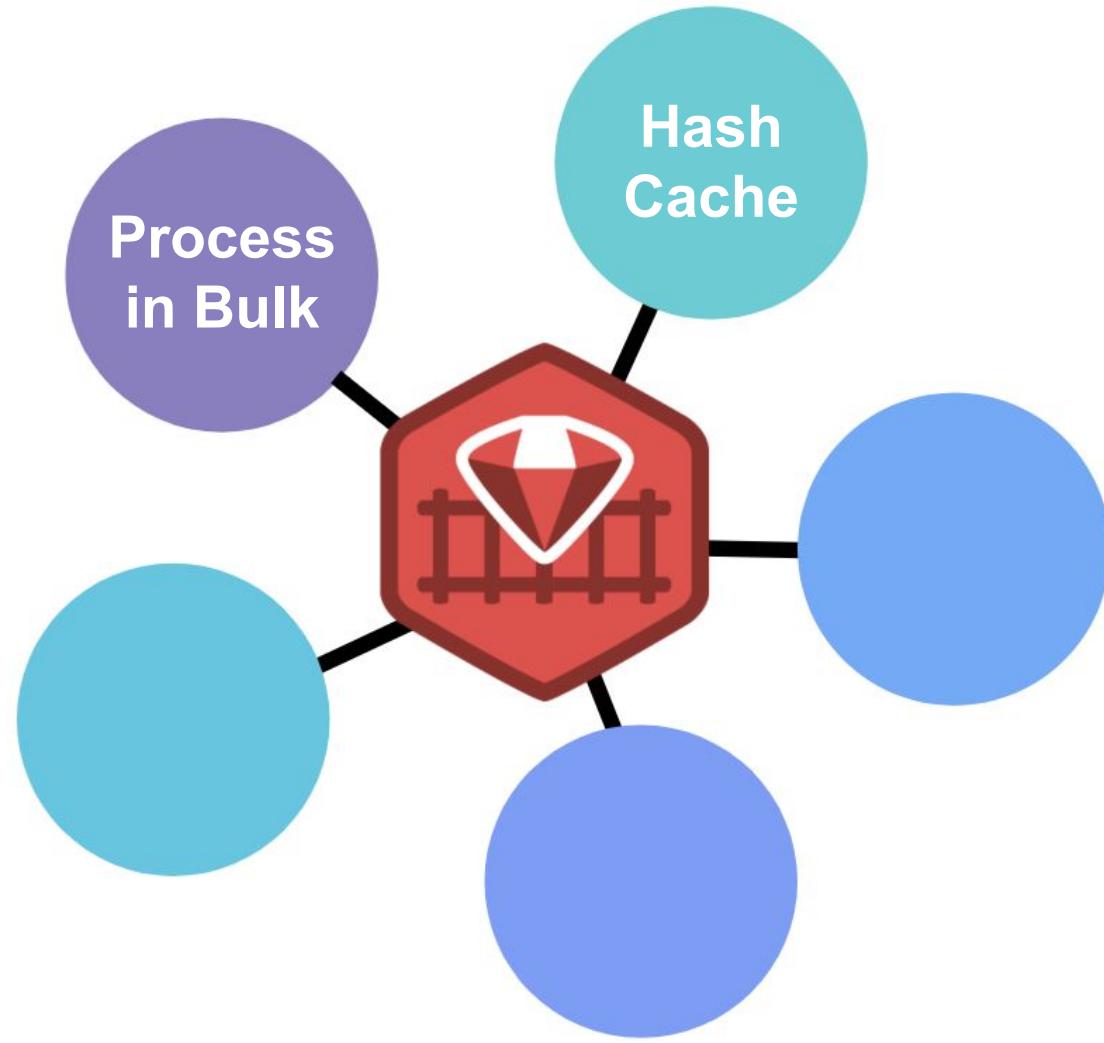
@molly\_struve

# Redis

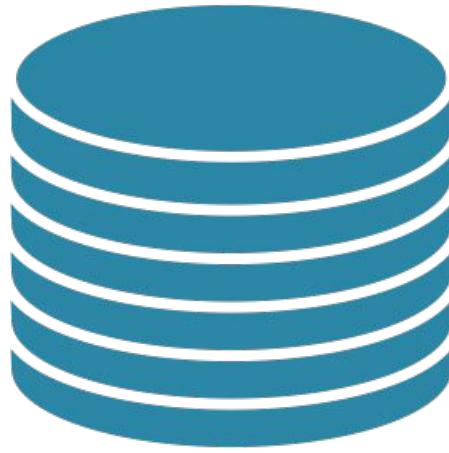
KENNA  
Security



@molly\_struve



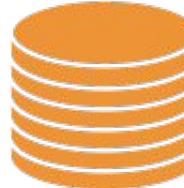
# Sharded Databases



**MASTER**  
Holds all application  
& reference data



CLIENT 1



CLIENT 2



CLIENT 3

# Octopus - Easy Database Sharding for ActiveRecord

[build](#) [passing](#) [CODE CLIMATE](#)

Octopus is a better way to do Database Sharding in ActiveRecord. Sharding allows multiple databases in the same rails application. While there are several projects that implement Sharding (e.g. DbCharmer, DataFabric, MultiDb), each project has its own limitations. The main goal of octopus project is to provide a better way of doing Database Sharding.

Asset.using(:shard\_1).find(1)

# Sharding Configuration Hash

```
{  
    'client_123' => 'shard_123',  
    'client_456' => 'shard_456',  
    'client_789' => 'shard_789'  
}
```

# Sharding Configuration Hash

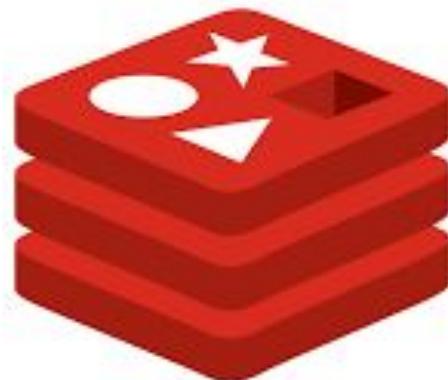
```
{  
    'client_123' => 'shard_123',  
    'client_456' => 'shard_456',  
    'client_789' => 'shard_789'  
}
```

}



sharding\_config\_hash[client\_id]

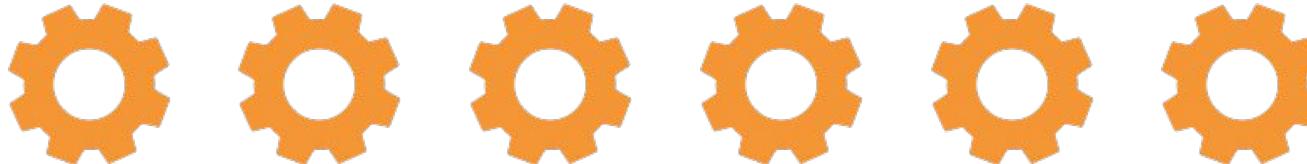




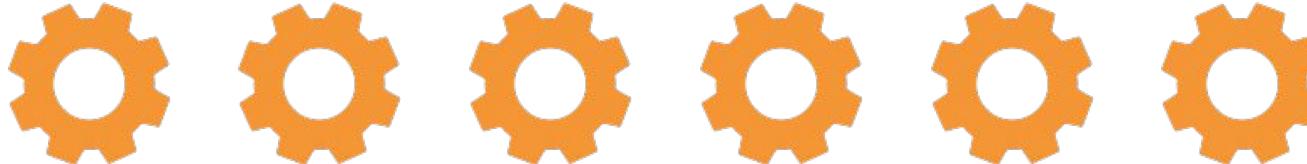
redis

# Sharding Configuration Size

20 bytes → 1kb → 13kb



**285 Workers**



**7.8 MB/second**

# ActiveRecord::Base.connection

```
(pry)> ActiveRecord::Base.connection
=> #<Octopus::Proxy:0x000055b38c697d10
@proxy_config= #<Octopus::ProxyConfig:0x000055b38c694ae8
```

## Octopus - Easy Database Sharding for ActiveRecord



CODE CLIMATE

Octopus is a better way to do Database Sharding in ActiveRecord. Sharding allows multiple databases in the same rails application. While there are several projects that implement Sharding (e.g. DbCharmer, DataFabric, MultiDb), each project has its own limitations. The main goal of octopus project is to provide a better way of doing Database Sharding.

@molly\_struve



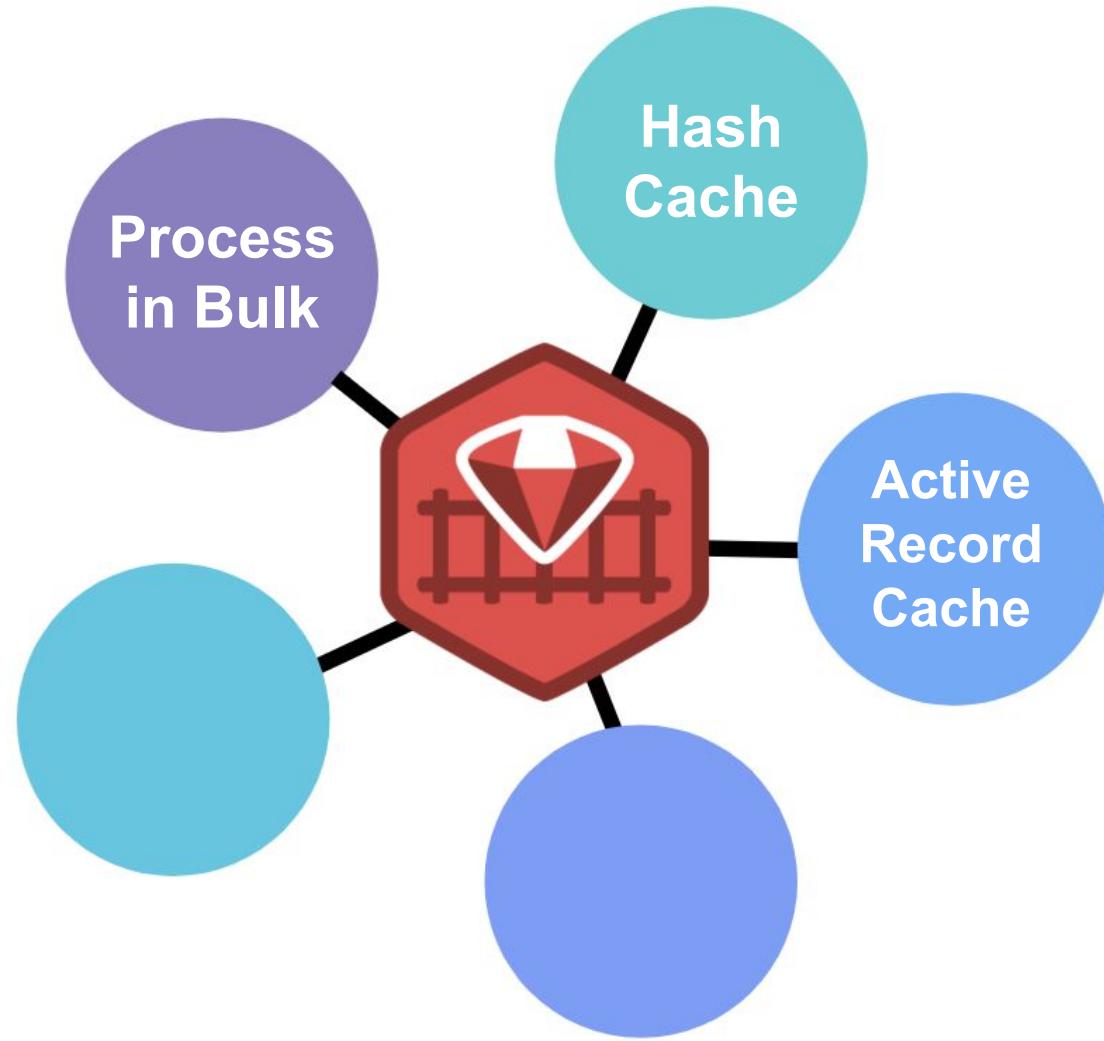
**SAY WHAT?**

```
module Octopus
  class Proxy
    attr_accessor :proxy_config
    delegate :current_shard, :current_shard=,
              :current_slave_group, :current_slave_group=,
              :shard_names, :shards_for_group, :shards, :sharded,
              :config, :initialize_shards, :shard_name, to: :proxy_config, prefix: false
  end
end
```



@molly\_struve

# Know your gems



# Avoid making datastore hits you don't need



@molly\_struve

```
User.where(:id => user_ids).each do |user|
  # Lots of user processing
end
```



```
(pry)> User.where(:id => [])
```

```
User Load (1.0ms) SELECT `users`.* FROM `users` WHERE 1=0  
=> []
```



@molly\_struve

```
return unless user_ids.any?
```

```
User.where(:id => user_ids).each do |user|  
  # Lots of user processing
```

```
end
```

```
(pry)> Benchmark.realtime do
> 10_000.times { User.where(:id => []) }
> end
=> 0.5508159045130014

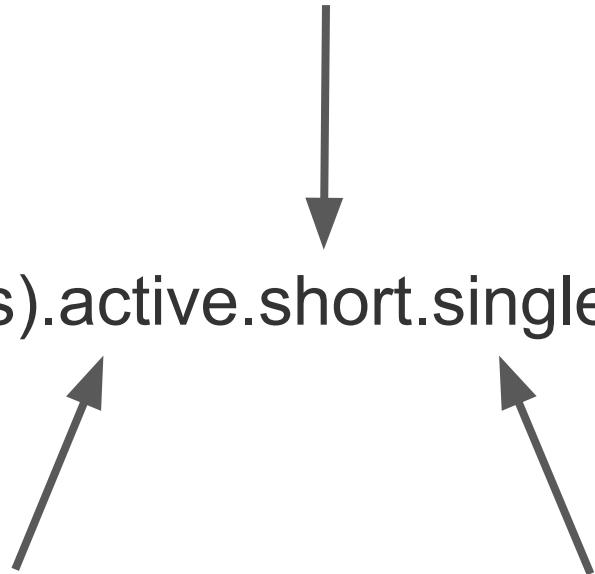
(pry)> Benchmark.realtime do
> 10_000.times do
>   next unless ids.any?
>   User.where(:id => [])
> end
> end
=> 0.0006368421018123627
```

```
(pry)> Benchmark.realtime do
>   10_000.times { User.where(:id => []) }
> end
=> 0.5508159045130014
```

Hitting the database is slow!

```
User.where(:id => user_ids).each do |user|
  # Lots of user processing
end
```

```
users = User.where(:id => user_ids).active.short.single
```



# .none



method

## none

---

Ruby on Rails latest stable (v4.2.7) - 0 notes - Class: ActiveRecord::QueryMethods

1.0.0 1.1.1 1.1.6 1.2.0 1.2.6 2.0.0 2.0.3 2.1.0 2.2.1 2.3.2 2.3.8 3.0.0 3.0.5 3.0.9 3.1.0 3.2.1 3.2.3 3.2.8 3.2.13 4.0.2 4.1.8 4.2.1 4.2.7

### none() public

Returns a chainable relation with zero records.

The returned relation implements the Null Object pattern. It is an object with defined null behavior and always returns an empty array of records without querying the database.

# .none in action...

```
(pry)> User.where(:id => []).active.tall.single
User Load (0.7ms) SELECT `users`.* FROM `users` WHERE 1=0 AND
`users`.`active` = 1 AND `users`.`short` = 0 AND `users`.`single` = 1
=> []

(pry)> User.none.active.tall.single
=> []
```



NEVER ASSUME, DEAR

@molly\_struve



Search Gems...

NEWS

GEMS

GUIDES

CONTRIBUTE

SIGN IN

SIGN UP

# Gems

---

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

---

DISPLAYING GEMS 1 - 30 OF 9619 IN TOTAL

@molly\_struve

# Logging

```
pry(main)> Rails.logger.level = 0
```

```
pry(main)> Search.connection.transport.logger = Logger.new(STDOUT)
```

```
$ redis-cli monitor > commands-redis-2018-10-01.txt
```

# Preventing useless datastore hits



**kenna** Home Dashboard Connectors Demo Inc.

## Windows Servers

October 04, 2016 View Top Fixes Explore

### Group Overview

Report Template: Set Ops Team

484 Assets
96,325 Vulnerabilities
17,870 Top Priority
5,170 Active Internet Breaches
15,825 Easily Exploitable
100 Zero Days



Highest Score

**700**

Lowest Score

**320**

5 MONTHS AGO

Vuln Density

**199**

Last Week

**700**

Last Month

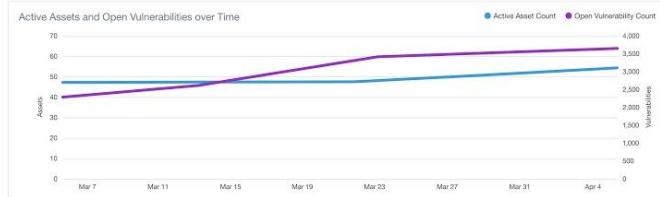
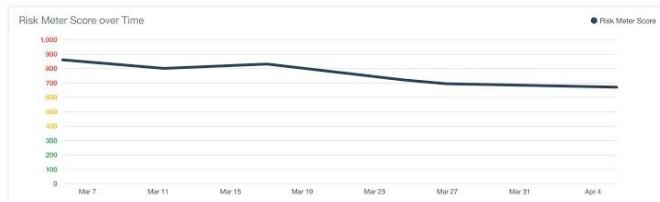
**320**

90 Days Ago

**320**

Group Created: November 23, 2015

### Risk Timelines



### Current Risk Information

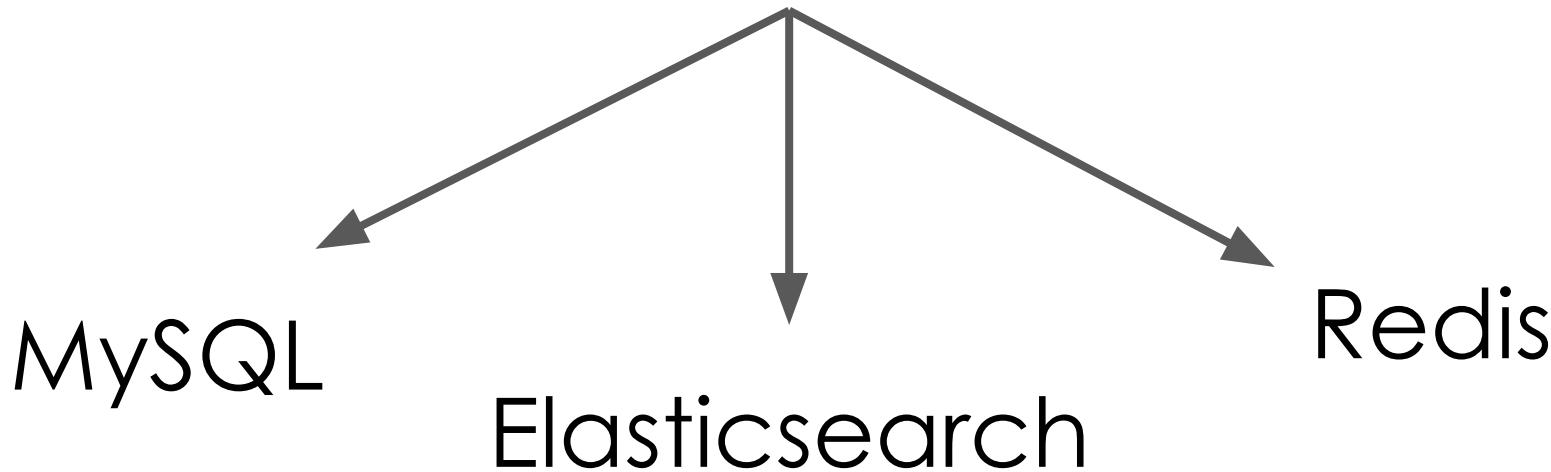
Open Vulnerabilities by Risk Level

High Risk	202
Medium Risk	920
Low Risk	2,591

Active Assets by Risk Level

High Risk	40
Medium Risk	19
Low Risk	5

# Report



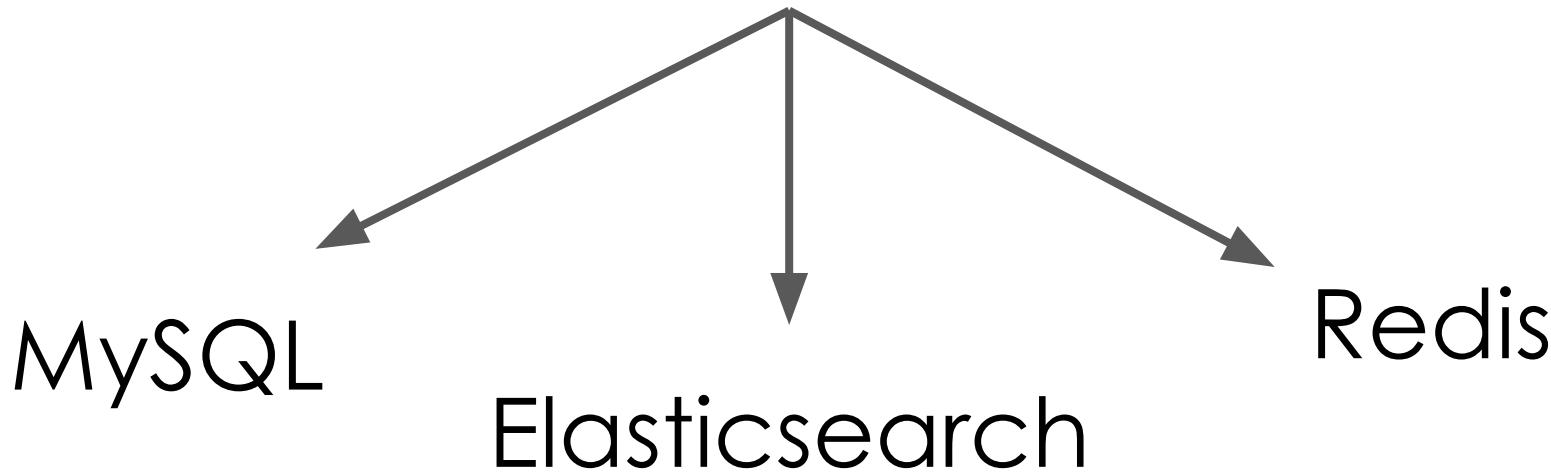
# Investigating Existing Reports

```
(pry)> Report.active.count  
=> 25842
```

```
(pry)> Report.average_asset_count  
=> 1657
```

```
(pry)> Report.zero_assets.count  
=> 10805
```

# Report



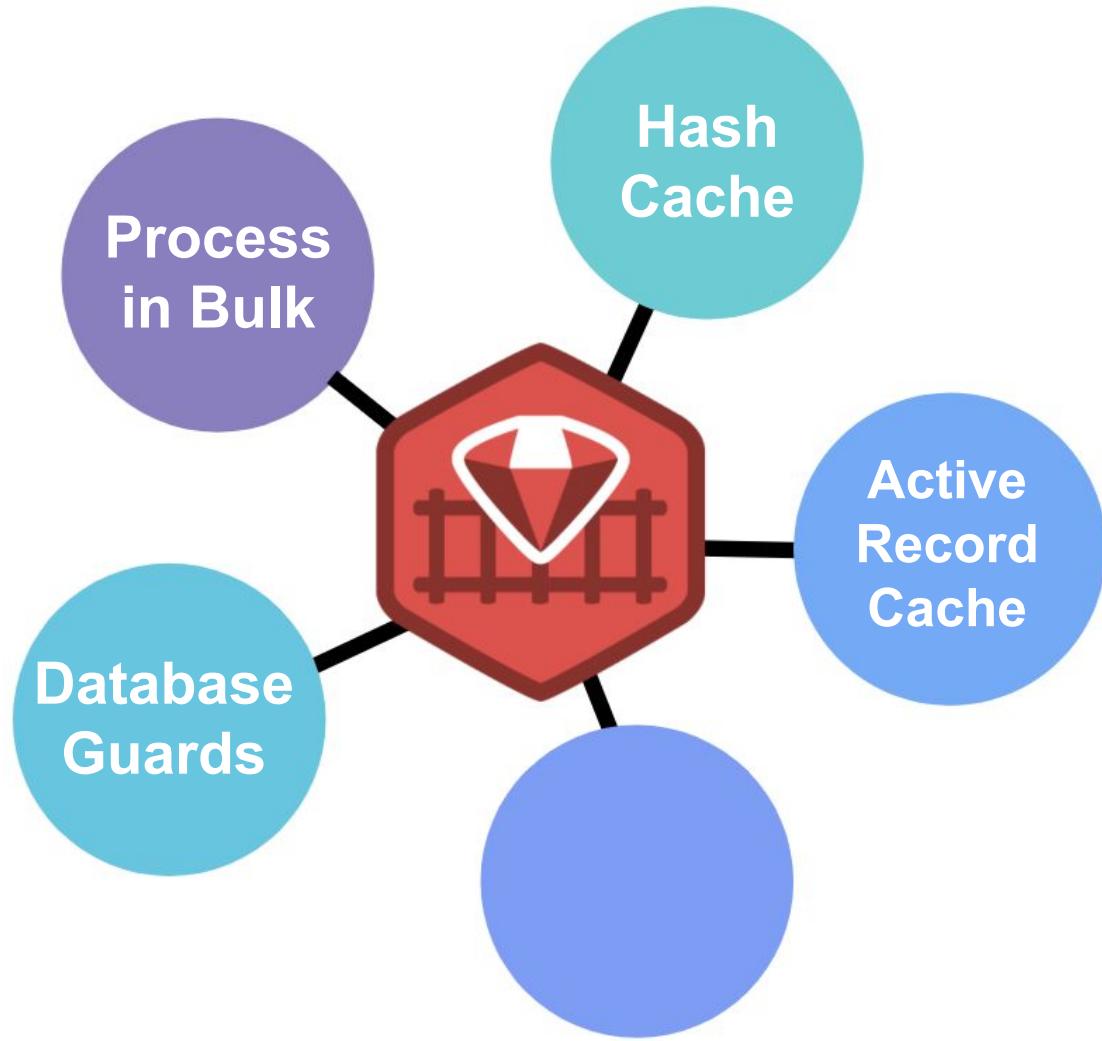


@molly\_struve

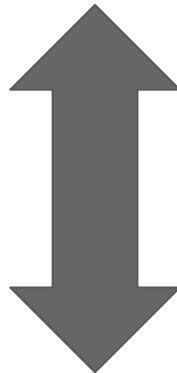
**return if report.asset\_count.zero?**

10+ hrs

3 hrs



# Resque Workers



## Redis



**45 workers**



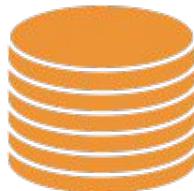
**45 workers**



**45 workers**



**70 workers**



**70 workers**



**70 workers**



## RDS CPU utilization



100

80

60

40

20

0

06:00

09:00

12:00

15:00

18:00



@molly\_struve



@molly\_struve



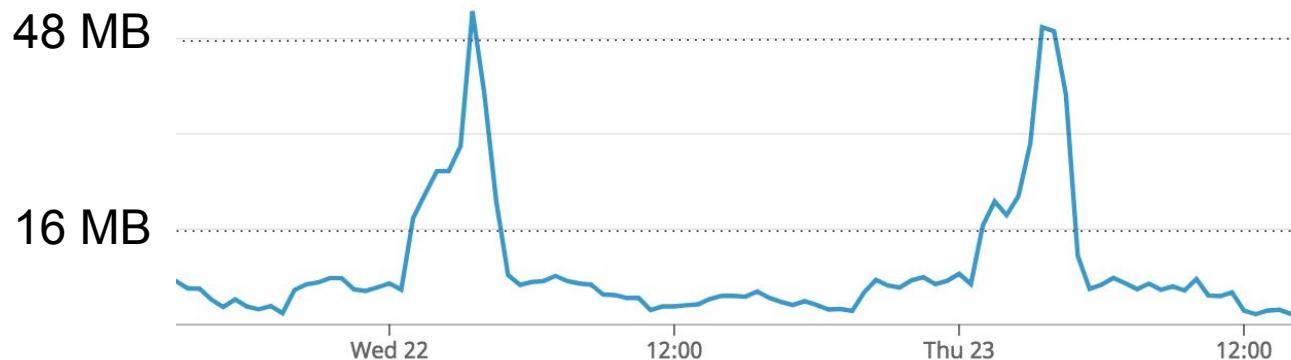
@molly\_struve



Datadog APP 3:05 AM

## Triggered: Redis: High Traffic

## High Redis Traffic Detected Call-Alert-Handling



@molly\_struve

# Redis Requests

200k



**70 workers**

## Redis::ConnectionError

A month ago - 09/05/18 @ 4:34 pm -05:00

### Summary

[Comments](#)[Backtrace](#)[Context](#)[Params](#)[Environment](#)[History](#)

#### Status

 **Unresolved**

We won't alert you again until the error is resolved.

#### Message

Redis::ConnectionError: Connection lost (ECONNRESET)

@molly\_struve



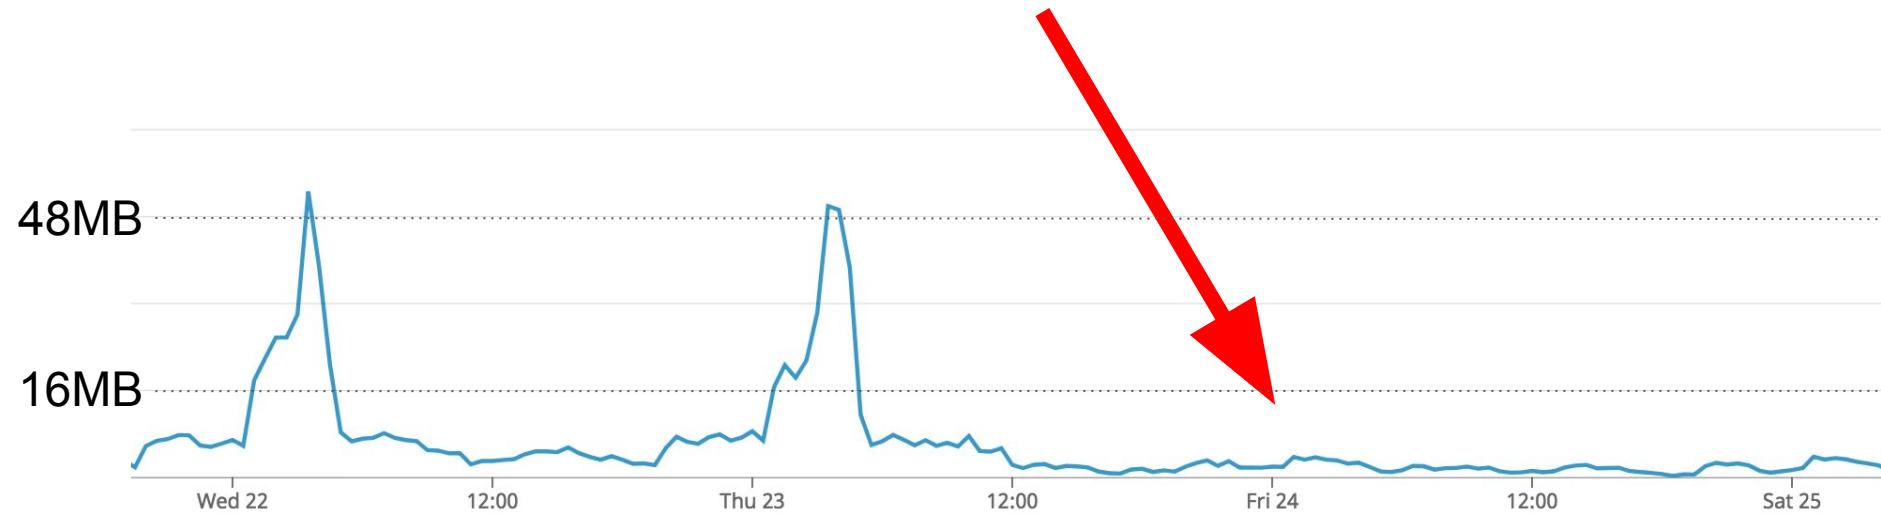


# Redis Requests



@molly\_struve

# Redis Network Traffic



@molly\_struve

## Redis::ConnectionError

A month ago - 09/05/18 @ 4:34 pm -05:00

### Summary

[Comments](#)[Backtrace](#)[Context](#)[Params](#)[Environment](#)[History](#)

#### Status

Resolved

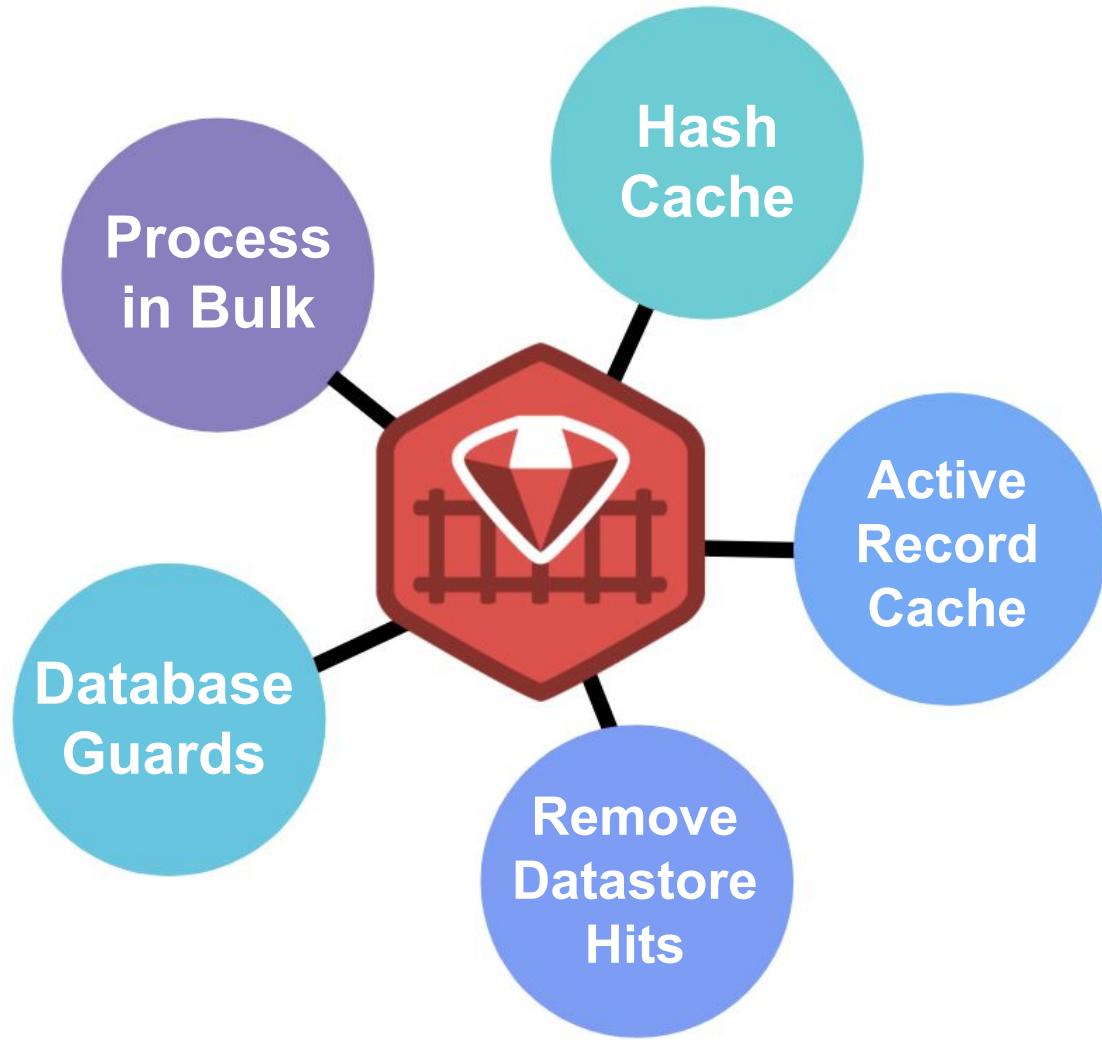


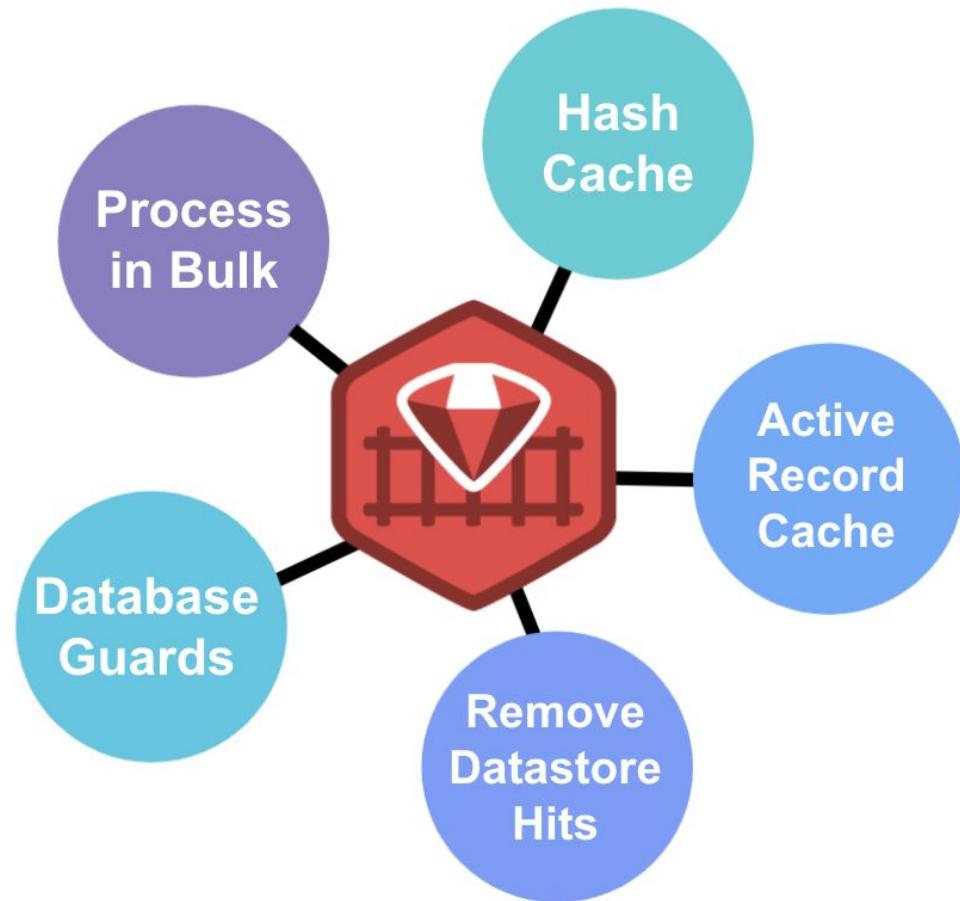
We won't alert you again until the error is resolved.

#### Message

Redis::ConnectionError: Connection lost (ECONNRESET)

@molly\_struve



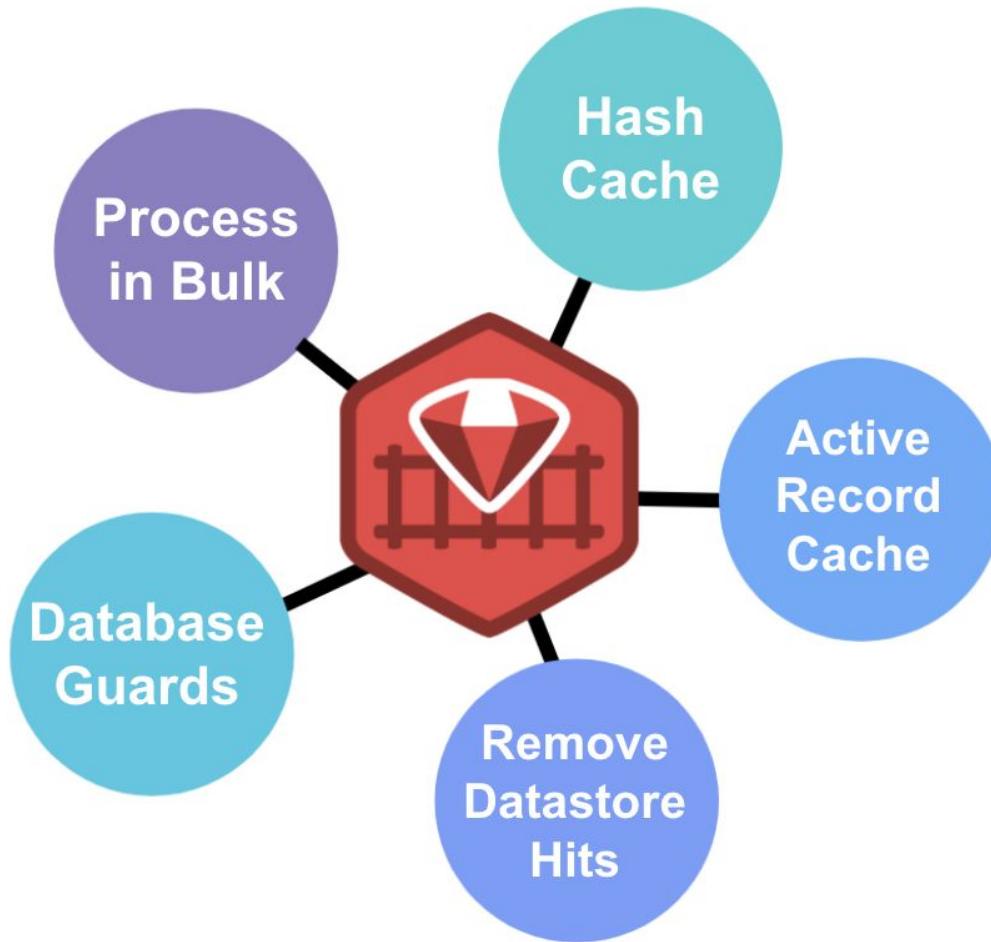


**Every  
datastore hit  
COUNTS**



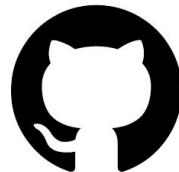
**My job here is done**

@molly\_struve



# Questions

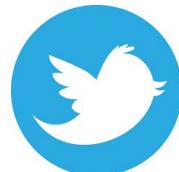
# Contact



<https://github.com/mstruve>



<https://www.linkedin.com/in/mollystruve/>



@molly\_struve



molly.struve@gmail.com

@molly\_struve